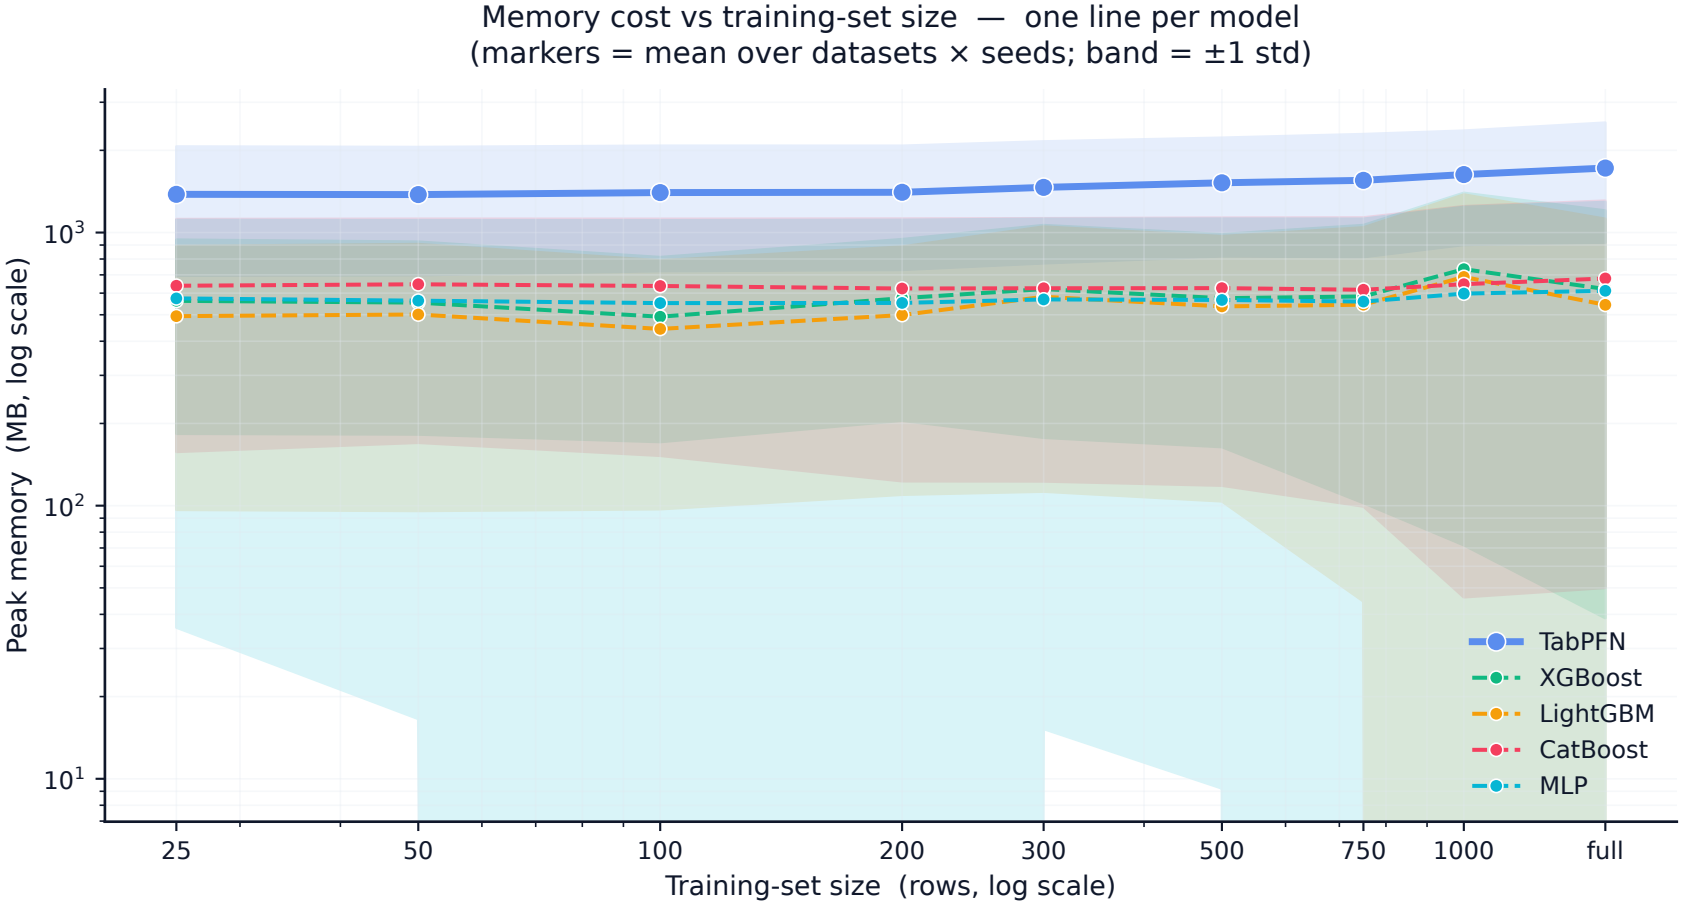
Memory cost vs training-set size — one line per model
(markers = mean over datasets × seeds; band = ±1 std)
Peak memory (MB, log scale)
$10^3$
$10^2$
$10^1$
25
50
100
200
300
500
750
1000
full
Training-set size (rows, log scale)
TabPFN
XGBoost
LightGBM
CatBoost
MLP